

Is Less Really More? Comparing the Climate and Productivity Impacts of a Shrinking Population*

Mark Budolfson
Univ. of Texas at Austin

Michael Geruso
Univ. of Texas at Austin
and NBER

Kevin Kuruc
Univ. of Texas at Austin

Dean Spears
Univ. of Texas at Austin
and IZA

Sangita Vyas
Hunter College

Abstract: A smaller human population would produce less carbon emissions, other things equal. This fact has led to the view that an important benefit of the ongoing, global decline in fertility will be reductions in long-run temperatures. Here we assess the magnitude and economic significance of this relationship. We find that it is quantitatively small and, therefore, insignificant relative to other well-documented effects of population growth, including non-rival innovation. This conclusion follows from a key fact of timing: Population sizes would respond to any growth rate changes begun today with a many-decades lag, by which point per capita emissions are projected to have significantly declined. Therefore, the additional warming from any plausibly-sized change in population growth would be small when compared against the resulting long-run productivity impacts. Moreover, even the *sign* of the population-warming relationship is ambiguous when we account for the possibility of aggregate net-negative emissions.

JEL Codes: Q54, J11, J13, O30, O40.

Keywords: Climate change; population dynamics; semi-endogenous growth; aging populations.

*All authors are affiliates of the Population Wellbeing Initiative at UT Austin. Geruso is at the Council of Economic Advisers. We thank Maya Eden, Frank Errickson, Jim Feyrer, Chad Jones, Peter Kruse-Andersen, Kyle Meng, Marta Prato, Noah Scovronick, Robert Socolow, Phil Trammel, Sam Trejo, David Weil, Anson Zhou, Stéphane Zuber, Valeria Zurla and participants at the NBER's conference on Declining Population Growth, the University of Texas' Workshop on Scale, the University of Oxford, and the White House Council of Economic Advisers for comments and conversations that have greatly improved the paper. Any remaining errors are our own. This research was supported by grant P2CHD042849 awarded to the Population Research Center at The University of Texas at Austin by the Eunice Kennedy Shriver National Institute of Child Health and Human Development and by the Musk Foundation. The content is solely the responsibility of the authors and does not necessarily represent the views of the White House or our funders.

1 Introduction

A smaller global population would produce fewer carbon emissions, holding all else fixed. This fact informs a widely held view that reductions in population growth can play an important role in mitigating the eventual damages from climate change (Casey and Galor, 2017; Bongaarts and O'Neill, 2018; Wynes and Nicholas, 2017). However, there are also well-documented productivity benefits of large and growing populations via market size effects and business dynamism (Karahan et al., 2019; Peters and Walsh, 2021; Hopenhayn et al., 2022), increased provision of non-rival goods (Jones, 2022; Peters, 2022), and a less retiree-heavy age structure (Vollrath, 2020; Fernández-Villaverde et al., 2023; Maestas et al., 2023).

Understanding which of these categories of impacts on living standards—climate or productivity—is likely to dominate in the near and distant future is important: Two-thirds of people now live in a country with below-replacement fertility and the global population is projected to begin shrinking within the next few decades (United Nations, 2022). Various evaluations of these competing channels indicate that the costs and benefits both have the potential to be large. For example: achieving 2°C (rather than 4°C) may avert more than 10 p.p. of GDP losses annually (see e.g., Burke et al., 2015; Howard and Sterner, 2017), and recent work by Jones (2022) suggests economic growth may end entirely under enduring population decline. But there has been remarkably little dialogue between, on the one hand, research using integrated assessment models to understand the climate costs of human activity and, on the other, work studying the relationship between population and economic growth. How these forces compare and what this implies for future living standards thus remains an open and critical question.

In this paper, we quantitatively assess the climate and productivity impacts of a shrinking population. We begin with a simple analytical model to clarify how changes in population growth (a flow) pass through to long-run cumulative emissions (a stock), arriving at a closed-form relationship between these terms. This solution demonstrates exactly how this link is mediated by the rate of (per capita) decarbonization and helps to develop intuitions for the dynamic interactions at the core of our analysis. Calibrating the relevant parameters to possibilities for decarbonization paths yields the novel result that the pass-through from population

growth to long-run warming is necessarily limited. This would not have been true a half century ago or more, but it is true from the vantage of a society that has already begun reducing per capita emissions. This insight—about the response of cumulative emissions to future population size—is critical for understanding why even moderate productivity losses caused by a shrinking world might outstrip the climate benefits of smaller population sizes.

The simple model illustrates the key mechanism over the long run, but does not deliver detailed quantitative comparisons or consider medium-run dynamics of the decades ahead. To generate these assessments, our main numerical exercise builds on William Nordhaus' DICE model (Nordhaus, 2017), the most widely-used and studied climate-economy model. We innovate on this model by incorporating two key ideas related to population growth: (i) that people are the source of non-rival innovations and ideas that propel economic growth, and (ii) that some consumers are too young or too old to also be workers, so that changes in the age structure of the population can affect per capita production for a given population size. The innovation effects of population are calibrated to recent estimates of the elasticity of productivity gains with respect to research inputs (Bloom et al., 2020). Within this framework, we contrast per capita output under two paths for future population: a shrinking population that is consistent with consensus projections and alternatively a stable-population future.

Consistent with the intuition from the analytic model, the result from the integrated climate assessment highlights the strikingly small temperature response to even large changes in future population size. Immediate and persistent increases to fertility, even to unrealistically high levels, produce a difference of 0.05°C by 2200 relative to status-quo projections of global depopulation in our quantitative model. This is on a pessimistic decarbonization trajectory where emissions remain positive for another century and eventual warming exceeds 4°C . This finding is not built-in as an assumption: We do not specify that the rate of decarbonization increases with population size through increased innovation or a planner's reoptimization. Instead, we conservatively and mechanically hold fixed the same time path of emissions-per-output-unit across the large population and small population scenarios.

The integrated model finds that increases in fertility, relative to the status quo population decline, would benefit living standards mainly through the contribution to productivity of increased non-rival innovation that a larger population can sustain. In contrast to conventional wisdom, but as anticipated by Weil (1999, 2023), the dependency ratio benefits are more limited, and are realized only after many decades. This is because the additional children in a higher-fertility scenario increase the near-term dependency ratio relative to the shrinking population. Overall, the additional climate costs resulting from a larger future population according to standard climate damage functions are small compared to the competing economic effects, all measured on a common scale of per capita economic output.

This conclusion follows from facts of timing. The cumulative concentration of atmospheric greenhouse gases—the driver of global warming—is a stock determined by the entire history of emissions. Moreover, while the annual flow of emissions can in principle adjust quickly, the instantaneous global population size is itself a slow-moving stock. Changes in population growth rates take many decades to generate significant changes in population size. So, the relationship between current population growth rates and long-run warming will be mediated by per capita emissions over the medium- and long-term. Even under pessimistic scenarios for decarbonization, only a small share of humanity’s cumulative emissions will occur after the many decades it takes to produce large population size changes. Scaling up or down this small remainder of emissions has only a small impact on the long-run stock of atmospheric GHGs. This implies that, in the long run, any form of gains associated with a larger future population need not be large to dominate the tiny difference in temperature generated by the larger future population.

Because our main results follow from these straightforward facts of timing, they are robust to a wide range of model specifications. For example, climate damage functions that translate temperature changes into harms are an area of uncertainty and active research in climate economics (Dell et al., 2012; Burke et al., 2015; Dietz et al., 2021; Cruz and Rossi-Hansberg, Forthcoming, etc.). Increasing the scale or shape of the climate damage function lowers long-run incomes (and hence increases the social cost of carbon) in any population scenario. But our core results are robust to several alternative (more pessimistic) specifications for the damage

function because any assumed damage function acts on a temperature difference too small to outpace the productivity benefits of a larger population. Likewise, because the small temperature difference between population scenarios is driven by a small difference in cumulative emissions, the findings are not sensitive to use of a more sophisticated climate and atmospheric representation, which translate from emissions to temperatures. The same is true for other model dimensions that we vary. The most critical assumption is over the pace of decarbonization. Along with our presentation of the main results, we demonstrate that, even in worst-case futures with more than 6°C of warming, the *incremental* effect of the larger population on warming and climate damages is small. This carries the important implication that climate policy and technology will mostly succeed or fail independent of population trajectories, and tells against describing population decline as even a useful complement to or component of climate policy.

These results do not imply that larger populations come at no cost. The time path of changes we estimate shows that there is an inter-generational trade-off: The net effect of the larger population scenario (which assumes exogenous near-term increases in fertility rates) leads to lower per capita output in the coming decades before the new, larger generations age into the workforce. Simply put, children are initially expensive.

Finally, and as an extension of the main exercises, we consider the scenario in which it is eventually possible to produce *negative* emissions at scale. This possibility is ruled out by construction in the main exercises, to isolate the effect of the additional emissions of a larger population while emissions remain positive. But already there exist activities that generate negative emissions. These include high-tech approaches (e.g., direct air capture) and low-tech approaches (e.g., reforestation). In a future with scalable versions of carbon-negative technologies, even the sign of the relationship between warming and population growth is ambiguous. A larger population produces more near-term emissions which warm the planet, but could later also produce more negative emissions into the indefinite future, which re-cool the planet. Considering net-negative aggregate emissions upsets (and even reverses) common intuitions about population and climate: For any fixed stock of cumulative historical emissions to date, capturing them from the atmosphere is a

global public good that can be more cheaply produced on a per capita basis by a larger population than a smaller one. Therefore, even a dogmatic planner with an objective function that only values achieving some long-run temperature target may prefer the population scenario with higher levels of near-term population growth depending on the expected arrival of negative emissions at scale.

This paper contributes to an active and influential debate on the climate effects of population dynamics. Our first contribution is demonstrating that changes to future population growth have, at worst, small negative impacts for climate outcomes. This finding is in direct contrast with the dominant view in academic and public discussions. Within economics, too, the view that we rebut is mainstream. For example, Kruse-Andersen (2023a) and Gerlagh et al. (2023) conclude that policies aimed at reducing population growth could have significant effects on carbon accumulation this century, in part because these authors assume that population sizes can change much more quickly than is realistic.¹ Galor (2022) hypothesizes that reductions in fertility can provide additional time to transition to clean energy sources, a conjecture based on earlier work in Casey and Galor (2017) that reports a headline finding that a simulated reduction in fertility can reduce annual emissions by 35% by 2100. Our paper clarifies why the Casey and Galor (2017) estimate can be correct and misleading for understanding the climate impacts of future population growth or decline: The climate impact of a 35% reduction in *the annual flow of* emissions by 2100 depends on the base emissions flow in that year and afterward—and is unlikely to be more than a small fraction of the accumulated stock of atmospheric greenhouse gasses. Understanding this nuance is especially important given that the pace of actual and projected decarbonization has substantially accelerated over just the past few years (Hausfather and Peters, 2020; Arkolakis and Walsh, 2023). And it underscores the usefulness of our analytic model, which highlights and separates the parameter describing per capita emissions as a time-dependent variable.

¹Kruse-Andersen (2023a) studies a high-population scenario with a working age population of just over 9B in 2100 (whereas the UN's 95% confidence interval for individuals aged 25-64 in 2100 extends to a maximum of about 5.8B, with potentially another 1.5B aged 15-24); Gerlagh et al. (2023) compare scenarios of total population that range from 8B to 12B by 2060 (the UN's 95% confidence interval for this year is 9.7B - 10.7B). Population forecasting is fairly mechanical on time-horizons of decades, so it would be very surprising to observe such large deviations from leading projections.

Our second contribution is to propose, for the first time, that a comparatively larger population could have positive climate impacts in the longer run. The sum total of all historical emissions presents a fixed-cost problem (i.e., not tied to current population size) from the perspective of a future generation employing a negative emissions technology. By relaxing the assumption that per capita emissions are bounded below by zero, we cast doubt on even the *sign* of the long-run population-warming relationship. This insight complements the few existing papers re-examining the importance of population growth for climate outcomes. Bretschger (2020) makes the theoretical point that the optimal path of fossil fuel use will be independent of population sizes, generating a null relationship in steady-state between population sizes and cumulative emissions. Our lesson is different, and extends to settings where the path of fossil fuel use is, in fact, sensitive to population sizes.²

Finally, we contribute to the long and ongoing debate on the relationship between population trajectories and per capita incomes by showing that common sources of economic benefits from population growth have much larger effects than (marginal) climate damages. This literature is vast, encompassing many forces.³ We do not aim to incorporate every channel by which increases in fertility could affect living standards. Instead, this paper casts new light on the most widely cited channel by which a larger population is hypothesized to reduce long-run living standards—climate harms. Our findings indicate that future exercises attempting to sort out the net effects of population growth can set aside the marginal climate damages as a second-order consideration.

²Alternatively, Budolfson and Spears (2021) shares our focus on the slowness of population changes, but makes a much weaker claim. They show that fertility reduction cannot be a “core” mitigation strategy by showing that ‘no climate policy’ scenarios still reach 6.4°C (as opposed to 7.1°C) under large declines in fertility rates. This leaves open the relevant questions of (1) what temperature reductions can be achieved on plausible decarbonization trajectories and (2) whether these reductions are economically significant. We show that they are much smaller than the 0.7°C difference found under their ‘no policy’ path and much less consequential for per capita income than other effects of population growth. See also Bradshaw and Brook (2014) and our own companion paper in the demography literature (Kuruc et al., 2023).

³Alongside the citations in the first paragraph, see recent work by Dasgupta et al. (2021), Henderson et al. (2023), Pindyck (2022) and others.

2 An analytical demonstration of dynamic interactions between emissions and population

This section proposes a tractable analytical framework that links near-term changes in population growth to long-run changes in atmospheric GHGs. Rather than quantify this effect in a rich numerical setting—the objective of Section 4—the goal here is to clarify the dynamic interactions that drive the main results in the simplest fashion. The key insight delivered by our model is that altering the population growth rate now will not meaningfully change the size of the population on a timeline relevant for impacting the stock of GHGs and long-run warming. This is because of the background rate at which per-capita emissions intensities are declining relative to any plausible change in the rate of global population growth.

Global warming is determined by the stock of atmospheric GHGs. Therefore, we focus directly on this GHG stock, which we denote as \mathcal{E} . For simplicity, we assume that this stock is equal to the time-independent sum of annual (flow) emissions, $E(t)$. That is, the amount of carbon in the atmosphere at any moment is the sum of all historical emissions.⁴ The present day is normalized to $t = 0$, such that $t < 0$ is the past.

$$\mathcal{E} = \int_{-\infty}^{\infty} E(t)dt \quad (1)$$

There is no time-subscript on \mathcal{E} because it is a representation of where the stock of atmospheric GHGs converges as the flow of emissions converges to zero. This is what we refer to as the “long-run” atmospheric stock of GHGs, which in turn leads to some fixed level of long-run warming.

The evolution of $E(t)$ can be decomposed into the evolution of population, $N(t)$, and per capita emissions, $\epsilon(t)$. The respective rates of change for population and per capita emissions are denoted g and d . In this simplified setting with constant rates of growth, we ignore cases where $g > d$ because this leads to an unbounded

⁴A more realistic model would include depreciation of this stock over time because CO₂ does eventually dissipate from the atmosphere. However, its atmospheric half-life is longer than one century, so for the purposes of this exercise a zero rate of depreciation is a reasonable simplification. We defer detailed climate modeling to our integrated assessment that begins in Section 3.

GHG stock and long-run warming.

$$E(t) = N(t) \times \epsilon(t) \quad (2)$$

$$N(t) = N_0 e^{gt} \quad (3)$$

$$\epsilon(t) = \epsilon_0 e^{-dt}, \text{ with } d > g \quad (4)$$

Throughout the paper we take the rate of per capita decarbonization as exogenous and independent of population sizes.⁵ We do this for two reasons. First, our main result is to demonstrate that the pass through from additional population growth to additional warming is small. Thus, for conservatism and transparency we do not build in the additional (and difficult to calibrate) channel by which a larger population could endogenously decarbonize faster than a smaller population via an innovation advantage. Our results would be strictly and straightforwardly stronger if such a channel exists. Second, and more substantively, an innovation advantage will be limited on the time horizon relevant for climate change for the same reason that emissions differences are limited: Population sizes change too slowly. Relatively small increases in research capacity many decades from now will not be a first-order factor in reducing per capita emissions this century (see also Kruse-Andersen, 2023a).

Equations 2 – 4 imply the following path of exponential decay for annual emissions.

$$E(t) = E_0 e^{-(d-g)t} \quad (5)$$

Equation 5 implies a simple analytical relationship between g, d and all future emissions. Denoting the pre-determined stock of atmospheric GHGs from historical emissions as \mathcal{E}^h , we have:

$$\begin{aligned} \mathcal{E} &= \mathcal{E}^h + \int_0^\infty E(t) dt \\ &= \mathcal{E}^h + \frac{E_0}{d - g}. \end{aligned} \quad (6)$$

⁵This has the implication that our model is precisely consistent with O'Neill et al.'s (2012) widely-cited estimate of a within-period unit elasticity: "CO₂ emissions from energy use respond almost proportionately to changes in population size."

Table 1: Marginal emissions from a 0.1 p.p. increase in population growth rates are small as a share of total emissions

Decarbonization rate ($d - g$)	Half-life of annual emissions	Cumulative future emissions	Marginal emissions	as % of \mathcal{E}^f	as % of \mathcal{E}
0.01	70 years	$100E_0$	$10E_0$	10%	6.7%
0.02	35 years	$50E_0$	$2.5E_0$	5%	2.5%
0.04	17 years	$25E_0$	$0.6E_0$	2.5%	0.5%

Notes: Increase in emissions from a 0.1 p.p. increase in population growth (from a baseline of 0.8% in 2022). The emissions increase depends on the baseline rate of decarbonization ($d - g$) in Equation 5—this determines the time left for population growth to influence cumulative emissions. The half-life column documents how long it takes for annual emissions to be half their current level for a given ($d - g$). Marginal emissions are in terms of current annual emissions, E_0 . The final two columns present this as a share of cumulative *future* emissions, \mathcal{E}^f , and cumulative *total* emissions (both past and future), \mathcal{E} . ($\mathcal{E} = \mathcal{E}^h + \mathcal{E}^f$) Even in the most pessimistic scenario where it takes 70 years for annual emissions to reach half their current level, the increase in emissions from a population growth increase are less than 7%. In more aggressive climate policy scenarios, the effect on total cumulative emissions is less than 3% of cumulative emissions.

Equation 6 makes clear that there are two forces limiting the pass-through from population growth rates to the stock of long-run atmospheric GHGs. First, there already exists a non-trivial concentration of GHGs, \mathcal{E}^h . This will be important for intuitions that rely on altering flows in order to affect the total stock. Second, population growth is only one of two channels which determine cumulative future emissions. Future emissions also depend on the evolution of annual per capita emissions, which as a matter of fact have evolved rapidly in recent decades and are forecast to continue to change as consequence of both industrial policy and technology. Formally, the derivative of \mathcal{E} with respect to the population growth rate g is:

$$\frac{\partial \mathcal{E}}{\partial g} = \frac{E_0}{(d - g)^2}. \quad (7)$$

It will be helpful to contextualize this with the simple numerical examples in Table 1. Starting with a central scenario where $(d - g) = .02$, the half-life of annual emissions is approximately 35 years. In other words, it will take until about 2060 for annual global emissions to be halfway to net-zero. Under this assumption, cumulative future emissions would be equal to $\frac{1}{0.02} = 50$ times current annual emissions, E_0 . For context, the existing stock of emissions, \mathcal{E}^h , is also equal to approximately $50E_0$.⁶

⁶Current total CO₂ emissions are about 38 billion tonnes, whereas cumulative historical emissions are about 1.7 trillion tonnes (see <https://ourworldindata.org/grapher/>)

So if $(d - g) = .02$, then $\mathcal{E} = 2\mathcal{E}^h$: Future cumulative emissions would equal past cumulative emissions, so long-run cumulative emissions would be twice historical cumulative emissions.⁷

Perhaps surprisingly, even under these assumptions, which imply significant future emissions, the marginal impact of future population growth on long run warming is small. Table 1 considers the effect of increasing g by 0.1 p.p. Such an increase would be large: it is well above the upper-bound of what estimates suggest could be accomplished with a one-time government benefit at birth equal to 10% of a household's annual income (Stone, 2020; Zhou, 2023).⁸ Indeed, it would be equivalent to rolling back more than 20% of the global fall in birth rates over the last two decades.⁹ According to Equation 7, this increase in g results in $2.5 \times E_0$ additional future emissions.¹⁰ In other words, this increase in g only increases cumulative future emissions, \mathcal{E}^f , by 5%. And what matters is \mathcal{E} , which will increase proportionally less than cumulative future emissions because of the pre-determined atmospheric stock, \mathcal{E}^h . As noted above, the pre-determined atmospheric stock is roughly equal to the quantity of future emissions under this assumption, implying that the increase in the overall stock is only 2.5%.

In a significantly more pessimistic scenario, where the decay rate on annual emissions is only 1% per year, this same 0.1 p.p. increase in population growth rates has a larger level effect. But this is on a base of cumulative emissions that is also larger. Overall, even under a very slow pace of decarbonization there is only a 6.7% increase in the long-run stock of GHGs. Conversely, in scenarios with faster rates of decarbonization there is no quantitatively significant relationship between population growth and the long-run stock of atmospheric GHGs.

cumulative-co-emissions).

⁷For further context, this central case would likely result in peak warming between 2.5-3°C. Current estimates of "committed warming" from existing atmospheric GHGs are between 1.2°-1.5° (Sherwood et al., 2022).

⁸The natural rate of population growth is approximately the birth rate minus the death rate. Holding fixed death rates, it takes a 0.1 p.p. increase in birth rates to raise the population growth rates by 0.1 p.p. The 2021 global birth rate was 1.69% (Our World in Data, 2023). Raising this to 1.79% represents a 5.9% increase in births. Stone (2020) and Zhou (2023) bound the fertility increase from these transfers below 4.1% and 3.65%, respectively.

⁹Between 2001-2021 the crude birth rate fell from 2.15% to 1.69%, or by 0.46 p.p. (Our World in Data, 2023). An increase of 0.1 p.p. undoes 22% of this decline.

¹⁰ $\frac{E_0}{(d-g)^2} = 2500E_0$ for a one-unit increase, or $2.5E_0$ for an increase of 0.001.

It is worth better understanding the intuition behind these small values as it is the key feature underlying the paper’s results. If we had naively assumed that N , rather than g , were immediately and permanently increased, the pass-through would be much larger. There would be a simple unit elasticity between cumulative future emissions and population increase because we have assumed a fixed path of per capita emissions. The difference arises from timing: an increase in g leads to a proportionally similar increase in N only after many decades. In our numerical example, increasing g by 0.1 p.p. generates a 10% increase in the population size only after 100 years.¹¹ Over the next century, emission intensities are declining. The population becomes significantly larger only *after* per capita emissions are much lower than their current level, limiting the ability for population size to influence the long-run stock of emissions.

This simple setting provides strong reason to conclude that the emissions, and therefore warming, effects of increases in population growth will be small. A richer quantitative model is necessary for demonstrating exactly how small, and how the implications of these small increases compare with the economic effects of population growth. That is the task of the next two sections. To preview: The marginal warming from changes to population growth is insignificant relative to the productivity benefits.

3 Integrated assessment model of population, the economy, and climate

The prior section demonstrated in a simple analytical model that the impact of population growth on warming could be small under various parameter assumptions, as long as humanity eventually decarbonizes. This section describes the quantitative model that assesses the same question using detailed population projections and a full integrated climate economy model. It also describes how the model can be used to study the tradeoffs between climate costs and productivity benefits.

¹¹For the time it takes to achieve a (counterfactual) 10% increase in the size of the population from a 1% increase in g , consider that $\frac{N'_t}{N_t} = \frac{N_0 e^{(x+0.001)t}}{N_0 e^{xt}} \Rightarrow 1.1 = e^{.001t} \Rightarrow t \approx \frac{.1}{.001} = 100$.

3.1 The DICE model

We start with DICE (Nordhaus, 2017), the most widely known and well-studied climate-economy model. DICE is not explicitly designed to study the implications of population growth, but is built on top of a neoclassical growth model in which labor contributes to output (and output determines capital accumulation, emissions, and so on). The model is therefore already constructed in a way that generates an emissions response to a change in the path of population. We build on this by introducing into the model the effects of population growth on productivity and the dependency ratio.

The four core features that we take directly from DICE are: (a) the neoclassical model of economic growth where labor, (accumulated) capital, and productivity determine total output, (b) a forecast for the time path of the emissions-intensity of output (absent climate policy), (c) a representation of how annual greenhouse gas emissions influence the atmospheric stock and how that stock in turn influences the global temperature over time, and (d) a damage function that translates temperature changes to losses of GDP.

Formally, gross output, Y^G , is defined by a standard Cobb-Douglas production function which includes capital, labor, and total factor productivity (TFP), A . Net output, Y^N , which can be used for consumption and investment, is what is left of gross output after emission-mitigation costs, Λ , are paid and climate damages, D , are suffered. Damages are represented as losses to GDP, but are calibrated to include the monetary value of non-market harms (e.g., health and mortality effects). D is assumed to increase quadratically in temperature T (above pre-industrial levels), although we explore alternative specifications. As we show in Section 4.2, our core results are robust to the substitution of other damage functions. The key components from DICE are thus:

$$Y_t^G = A_t K_t^\gamma L_t^{1-\gamma} \quad (8)$$

$$Y_t^N = (1 - \Lambda_t)(1 - D_t)Y_t^G \quad (9)$$

$$D_t = \psi_1 T_t + \psi_2 T_t^2 \quad (10)$$

Industrial emissions, E_t , are a function of gross output, determined by the emissions-

intensity of production, σ_t . These emissions can be abated at rate μ_t , the standard climate policy variable in analyses using DICE, which determines cost Λ in Equation 9:

$$E_t = (1 - \mu_t)\sigma_t Y_t^G \quad (11)$$

$$\Lambda_t = \theta_1 \mu_t^{\theta_2} \quad (12)$$

For brevity, we omit description of the climate and atmospheric modules, which map the history of E_t (and non-industrial emissions) to T_t . They are discussed in detail elsewhere (see e.g., Nordhaus, 2017), and are not directly relevant for the economic lessons of this paper.¹²

3.2 Modifications

Our first major modification of DICE is to incorporate the innovation benefits of population. The original DICE model assumes that total factor productivity increases are exogenous. Instead, following the semi-endogenous growth literature (Jones, 1995, 2022), we allow for resources—namely, people—to contribute to economic growth. These models build on the insight by Romer (1990) that larger economies produce more non-rival goods in aggregate (such as ideas and innovations), which increase per capita productivity.

Specifically, we employ a canonical semi-endogenous growth equation (Jones, 1995):

$$g_{A,t} = \frac{\Delta A_t}{A_t} = \alpha_t L_t^\lambda A_t^{-\beta} \quad (13)$$

The rate of increase of A is increasing in the size of the labor force, L (not the population). Innovation and progress comes from economic activity—either through learning-by-doing or explicit research efforts—which scales with L . α_t is a scaling factor between the labor force and the production of ideas, determined by the share of the labor force participating in idea production as well as the productivity of this sector. λ allows for intra-period increasing or decreasing returns to research

¹²Furthermore, we verify that the atmospheric details of the model are not consequential for the main results: Robustness exercises that replace the DICE atmospheric module with a focal alternative representation of this process (National Academies, 2017) generate nearly identical results.

effort. $\beta > 0$ allows for the possibility that there are dynamic diminishing returns to knowledge accumulation. If proportional increases in productivity become more difficult to achieve as productivity increases, that would be reflected by a large value for β . As in our simple model of Section 2, there is no link between TFP, A , and the emissions intensity of output, σ . For conservatism, both population scenarios that we study here (and describe further below) face the same exogenous path of σ_t .

Our second major modification of DICE is to include dependency ratio effects. Because DICE was not designed to explicitly study changes in population growth rates, the standard model assumes that workers scale linearly with the population and therefore omits any distinction between workers and people. We decouple the total population from the work force based on the age structure in each period of the respective population scenarios described in Section 3.4. In Equation 8, L is the working-age population, which is not equal to the total consuming population, N , in our implementation. Accordingly, the working-age population ratio is $\frac{L}{N}$ and the dependency ratio is $1 - \frac{L}{N}$.¹³ This accounting is important because, for example, a shrinking population is an aging population, with a greater share of retirees.

As a final minor modification to DICE, we endogenize the emissions from changes in land use (E_{land} , e.g., from deforestation). These are exogenous in DICE, corresponding to its single exogenous path of population. We assume a unit-elasticity between population and land-use emissions, such that for population path m in time t :

$$E_{land,m,t} = \frac{N_{m,t}}{N_{DICE,t}} \times E_{land,DICE,t}. \quad (14)$$

If the population is $x\%$ larger in time t than it is in DICE for that period, land-use emissions will also be $x\%$ larger than in DICE for that period.

We do not explicitly model fertility decisions, or any other channel by which populations endogenously evolve. The analysis aims to understand the consequences of a shrinking population versus a stable population, not which trajectories are more or less likely. Section 3.4 discusses the construction of these two exogenous

¹³Modifying the labor input in this way implies an immediate and permanent decrease in L relative to DICE, where every person is assumed to be in the labor force. To avoid mechanically reducing total production from this redefinition, we add a constant scalar on labor productivity equal to $\frac{N_{2020}}{L_{2020}}$ to replicate year 2020 output.

population paths.

Parameter values for these DICE modules described above are taken directly from DICE where possible.¹⁴ For the semi-endogenous growth parameters we introduce into the model, we use $\lambda = 1$ and $\beta = 3.1$, based on Bloom et al. (2020). We calibrate α_t to exactly match DICE’s exogenous path of TFP growth when the population trajectory from DICE is inputted. Our goal is to replicate the baseline DICE model as closely as possible in an effort to isolate the effects of population, not any modifications of DICE itself. We verify that when the DICE population is read into our model, the exact output from DICE is obtained for all variables (see Appendix Figure A1).

In summary, the modifications to DICE are as follows: (i) Technological progress increases in population size based on the endogenous growth literature; (ii) the distinction between total population and labor is explicitly represented, such that an economy with more children or retirees has lower GDP per capita, other things equal; and (iii) emissions from deforestation and other sources of land use scale with population. Alternative model specifications presented in Section 4.2 additionally modify the climate damages in DICE, replace DICE’s climate module with the Finite Amplitude Impulse Response (FaIR) climate module in line with recommendations from the National Academies (2017), and increase the emissions impact of population. These analyses demonstrate that our results are robust to these modeling choices.

3.3 Decarbonization Paths

As made clear in Section 2, to study the climate costs of population paths, a stance must be taken on a decarbonization path. Advances in renewables technology and the implementation of (some) mitigation-inducing policy has rendered common “business as usual” emissions paths pessimistic relative to updated estimates of the world’s likely emissions and warming trajectory (Hausfather and Peters, 2020; Ou

¹⁴The version we modify is DICE2016—the latest version available at the time of research—which is publicly available on Nordhaus’ website (<https://williamnordhaus.com/dicerice-models>) and has been translated to other software and coding languages that we build from (see <https://www.mimiframework.org/>). Since performing this research, DICE2023 has been released. Most of the modifications align with modifications we had independently made in the robustness checks of Section 4.2.

et al., 2021; Arkolakis and Walsh, 2023). In our baseline case, we assume a path of mitigation rates calibrated to global emissions in 2030, 2050 and 2100 under the current policy trajectory estimate in Ou et al. (2021) (see Appendix Figure A2). This assumed “current policy” trajectory exhibits reductions of (net) emissions by the end of this century, but too slowly to meet international climate goals (see Figure 2). For conservatism, we also consider a “low ambition” policy environment, which yields end-of-century warming similar to common worst-case climate scenarios (Figure 4). In Figure 5 we additionally consider an alternative climate policy path that is much more ambitious than the baseline. (See also Section 4.2.)

Each comparative analysis between the two population scenarios holds decarbonization rates fixed, so the only source of emission differences will be the first-order effect that economic activity (including land use) is increasing in the size of the population. As noted in Section 2, alternatively allowing for endogenous channels by which the larger population can offset some of this first-order difference would only strengthen our main takeaway.¹⁵


3.4 Population paths: *Depopulation* versus *Stabilization*

We compare two paths for the long-run global population, which we call *Depopulation* and *Stabilization* for ease of reference. They are plotted in Figure 1. *Depopulation* represents demographers’ central, consensus projection of the demographic future (United Nations, 2022; Basten et al., 2013; Raftery and Ševčíková, 2023): Fertility rates worldwide will converge to below-replacement levels (i.e., lower than 2.1 births per woman) and global population growth will become negative later this century. Quantitatively, we use and extend cohort-component based population projections by the United Nations (UN) World Population Prospects (United Nations, 2022). These projections disaggregate the population into 5-year age bins, where age-specific fertility estimates are applied to each group to determine the size of the following periods’ cohort of newborns. This level of detail for the age-pyramids, and age-profile of reproductive rates, provides an accurate representation of the lag between changes in fertility rates and changes in the total population size.

¹⁵Aside from the possibility of increased progress on clean energy technology, we also have in mind here endogenous policy that responds to the anticipated increase in cumulative emissions and warming (Bretschger, 2020).

Our implementation of the *Depopulation* path replicates the UN Medium projection exactly until 2100, when that projection ends. After 2100 we mechanically project continued (negative) population growth, guided by Basten et al. (2013) and Spears et al. (2023), the latter of which is a companion paper detailing the long-term decline scenario studied here. In this population path, the global total fertility rate converges to 1.66 births per woman, the 2021 TFR in the United States, through the end of the model span.

Figure 1: Two Population Paths: *Depopulation* and *Stabilization*



Notes: *Depopulation* and *Stabilization* population paths (inputs to the evaluation in later sections) are derived from United Nations (UN) World Population Prospects 2022 projections. UN projections are available until 2100. *Depopulation* follows UN Medium until then, after which it is extended to match demographic facts for low-fertility populations (United Nations, 2022; Basten et al., 2013). *Stabilization* avoids population declines by bounding fertility rates from below at the replacement rate. Countries are immediately increased to this level if they are currently below replacement; if they are above replacement in 2020-2025, they follow their UN-Medium projected decline until reaching replacement rate.

In the *Stabilization* path, negative population growth is avoided and the population stabilizes. This purely hypothetical scenario is constructed by combining two existing UN scenarios for age-specific fertility rates: the Instant Replacement variant for High-income, Upper-middle-income, and Lower-middle-income countries, according to World Bank income groups, and the Medium variant for Low-income,

and No-income-group-available countries (so these latter two country groups have the same path to 2100 in *Stabilization* and *Depopulation*).¹⁶ The Instant Replacement variant is one in which fertility rates immediately increase to the replacement rate (a TFR of about 2.1). After 2100, when the UN projections end, fertility rates in richer countries remain at replacement rate, and fertility rates in Low-income and No-income-group countries converge to replacement if they have not already reached that level. In this scenario, the long-run global population stabilizes at about 13 billion people.¹⁷ The age pyramid stabilizes with about half of the population in working ages at any give time in later centuries. This is approximately the working-age proportion of the population that the UN projects for 2100 in High-income countries.

Figure 1 demonstrates the overall population pathways and highlights the distinction between stocks and flows that are important for understanding our results. Despite the immediate jump in fertility rates in the *Stabilization* path, and a population that is nearly twice as large by 2200, it takes until the end of this century for the difference in population size between the two paths to exceed just 10%. The main results follow from this (lack of) difference in population size during the time in which the world is expected to make progress towards net-zero emissions.

4 Results: Differences in emissions, temperature and per capita output from simulated differences in population growth

This section discusses our main results, first showing that the integrated assessment model of Section 3 generates emissions and temperature differences between the two population paths that are quantitatively small. It then assesses the economic impacts of productivity differences across the two population paths, which are large

¹⁶The distinction between countries by income level is purely functional for the construction of the *Stabilization* path. Country income level is correlated with fertility. The strategy described here brings below-replacement fertility countries (which happen to also be richer) up to replacement level, and allows above-replacement fertility countries (which happen to be poorer) to continue their UN-projected decline until 2100. In DICE, we input paths for the *global* labor force and population size, and do not disaggregate by country.

¹⁷Our results are not contingent on these particular population paths. The productivity benefits of population far outweigh the climate damages even in a comparison between *Depopulation* and a much higher population projection, one which is well beyond the upper bound of the UN's 95% prediction interval for the demographic future (see Appendix Figure A4).

in comparison. The section concludes by demonstrating the robustness of these results to a host of alternative parameterizations and specifications.

4.1 Main Results


Consider first Figure 2, which depicts the emissions and climate impacts of replacement-rate fertility under current policy estimates. Recall that *policy* here means only the exogenous time path of per capita annual emissions. The temperature in the two population scenarios is nearly indistinguishable on the scale plotted, following directly from the slow speed of population change relative to the decline in emissions intensities that we have already stressed. *Depopulation* reaches 4.17°C by 2200; *Stabilization* reaches 4.22°C .¹⁸ While these are warming levels that are likely to result in substantial climate damages, the point we make here is that the expected damages are effectively invariant to the path of population. The difference between them (0.05°C) is small. This 1.2% difference in warming would, under a quadratic damage function, lead to a difference of less than 3% in annual climate damages in 2200. By that same point, the difference in population would be nearly 6 billion people, an 87% difference.

Because the climate costs of a larger population are so small, even very modest benefits of population arising from endogenous innovation or dependency ratio improvements can dominate the harms. There are two reasons for this. First, an 87% increase in population by 2200 would require a long-run population-productivity elasticity of only about 0.03 to generate a similar 3% increase in the level of productivity. Second, climate damages are only a fraction of GDP. A 3% increase in climate damages that are, for example, 10% the size of total output produces just a 0.3% decline in output available for consumption and investment. Increases in total factor productivity, or the labor force, would instead scale total output. So in fact, the long-run population-productivity elasticity would not even need to reach 0.01 to exceed climate damages.

Evidence from the literature on population and productivity (Bloom et al., 2020; Peters, 2022; Ekerdt and Wu, 2023; Kruse-Andersen, 2023b) is consistent with a

¹⁸These temperature changes we find under a current policy scenario are large, compared to the focal 1.5°C target outlined in the Paris Agreement. That is because the current policy scenario we evaluate is not optimistic.

Figure 2: Long-run warming is nearly indistinguishable between population scenarios



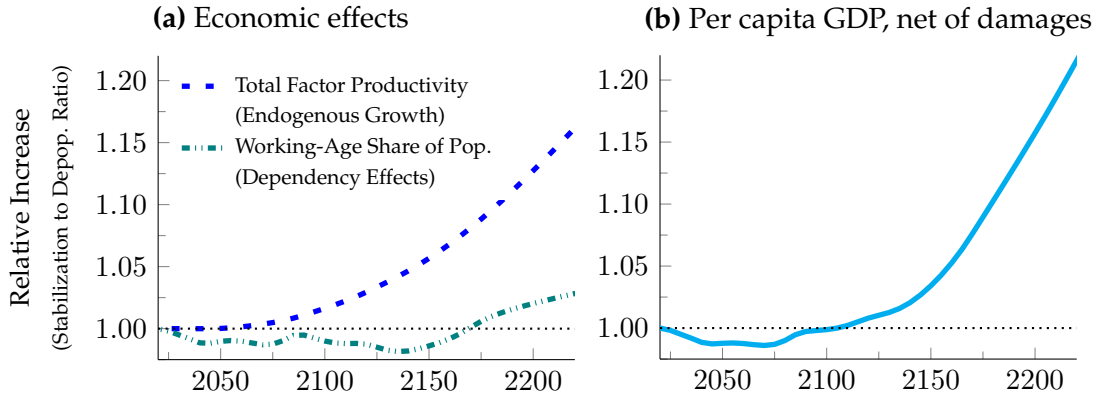
Notes: Emissions (left-axis) and temperature above pre-industrial (right axis) in each of the two population paths described in Section 3.4 when read through the climate economy model. Annual emissions in each year approximately scale with the population—for example, emissions are approximately 10% higher in 2100, just as population is—but are a small share of remaining cumulative emissions by the time population change becomes significant.

population-productivity elasticity of at least 0.3—about thirty times larger than the minimum needed for productivity gains to exceed climate damages in our baseline scenario in Figure 2. Figure 3a documents the relative increases in TFP and the share of the population in working ages for our larger population scenario. By 2100, TFP is roughly 2% larger in *Stabilization*. This relatively small effect by the end of the century is because TFP is also a cumulative stock, and population size increases slowly. By 2200 TFP is more than 10% larger in *Stabilization* relative to *Depopulation*.

The dependency ratio impacts are more complex, tracing out non-monotonic effects. Initially, and for a prolonged period, the additional children worsen the dependency ratio, leading to lower GDP per capita, a finding anticipated by Marois et al. (2021), Vollrath (2020), Weil (1999, 2023) and others. Over the long run, however, the dependency ratio improves. *Stabilization* eventually generates an age profile advantage in which about 2.5 percentage points (about 5%) more of the population are workers, relative to *Depopulation*.¹⁹

¹⁹This analysis is not designed to make normative claims about changes in population growth.

Figure 3: Economic effects of population stabilization exceed climate costs



Notes: Panel (a) plots the increases in total factor productivity and working-age share under *Stabilization* relative to *Depopulation* as ratios. These are invariant to the assumed climate policy. Panel (d) plots the increase in average living standards (measured on scale of per capita income, net of climate damages) in *Stabilization* relative to *Depopulation*. This main outcome holds under a wide range of variations on baseline assumptions (see Figure 5).

4.2 Robustness

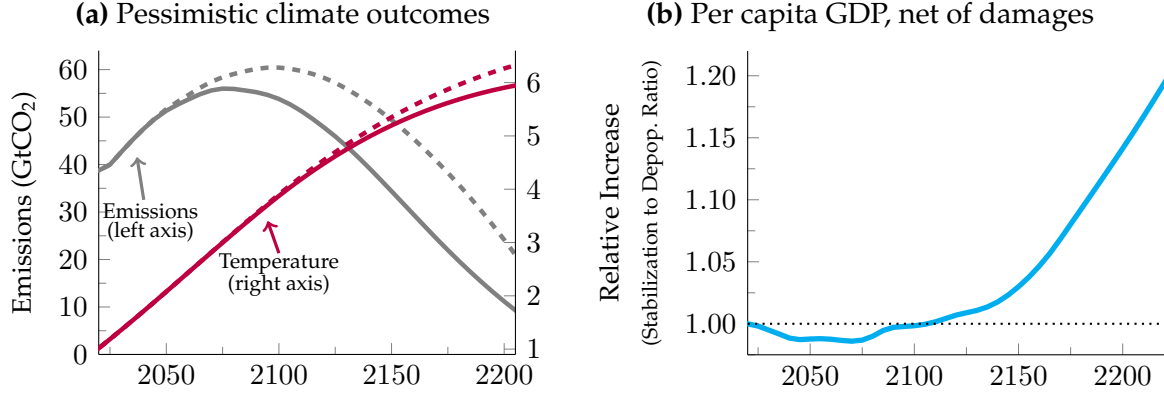
The magnitude of population's impact on climate is conditional on assumptions about decarbonization progress. Realistically accounting for the dynamic interaction between demographic change and decarbonization progress is a contribution of this paper relative to past work. But it also implies that if per capita emissions remain significant into the far future, the warming effects of near-term population growth could become large. To provide a sense for whether plausible, but extremely pessimistic, decarbonization paths could reverse our qualitative conclusions, Figure 4 considers a pathway that is in line with standard worst case scenarios (Hausfather and Peters, 2020; Ou et al., 2021). In this future, warming is approximately 4°C by the end of this century and 6°C by 2200, even with long-term population decline.

We find that in this extreme warming scenario the *marginal* contribution of population growth to warming remains small: 6.26°C versus 5.90°C (Figure 4a). For the same reasons as above, this 0.36° difference continues to imply a small impact on per capita GDP-equivalent damages, relative to the innovation benefits. Panel (b) makes this explicit by plotting the net effect of larger populations in this pessimistic

Therefore, we do not attempt to take a discounted sum of changes to per capita output over time. The positive research question that motivates the analysis is whether low fertility in 2023 and beyond can have an economically significant effect on long-run climate outcomes.

decarbonization scenario. It is positive in the long-run and qualitative similar to the results in Figure 3b.

Figure 4: Main results are robust to using more pessimistic decarbonization paths



Notes: These figures plot identical exercises to those that generated Figure 2 and Figure 3b, respectively, with a more pessimistic assumption for the time it takes to decarbonize.

Alongside varying the emissions pathways, we gauge sensitivity to alternative implementations of the endogenous growth channel, and to addressing several well-studied limitations of DICE. Figure 5 presents the results of 192 robustness checks, each from a different set of modifications to the baseline model. Scenarios in Figure 5 span from ambitious futures in which temperature change comes close to meeting international targets and global living standards grow four-fold by 2100, to scenarios with end-of-century temperature change near 4°C and in which living standards *fall* over the following century. In all cases the *additional* warming caused by larger populations remains small and economically insignificant.

These 192 modifications come from interacting changes to climate policy, the climate and atmospheric module, the climate damage function, parameter values in the semi-endogenous growth equation, the driver of increased TFP growth in the endogenous growth equation, and the elasticity of emissions with respect to policy.

$$\underbrace{\text{Climate Policy}}_3 \times \underbrace{\text{Climate Representation}}_2 \times \underbrace{\text{Climate Damages}}_4 \times \underbrace{\text{Population} \rightarrow \text{TFP Pass Through}}_2 \\
 \times \underbrace{\text{Source of TFP Growth}}_2 \times \underbrace{\text{Population Emissions Elasticity}}_2 = 192$$

Each modification is described below.

Climate Policy. Three climate policy scenarios are considered. The first two have already been detailed as the baseline and pessimistic policies considered in Figures 2 and 4. For completeness we also include a more optimistic path, one in which 2°C end-of-century warming is only just breached. In this path, temperature differences between *Stabilization* and *Depopulation* are zero to more than three decimal places because annual emissions fall to zero well before significant population change occurs. (See Appendix Figure A3 for a detailed depiction of this scenario and the difference in warming between population scenarios.)

Climate Representation. The DICE climate representation was designed to integrate simply within a macroeconomic model. In recent years there have been numerous attempts to produce more realistic, but still tractable, climate representations. The Finite Amplitude Impulse Response (FaIR) is one such model that has been recommended in a National Academies’ report on better practices in integrated assessment modeling (National Academies, 2017).^{20,21}

Because FaIR may be of special interest to readers in the IAM community, we additionally replicate Figure 3 in Appendix Figure A5. FaIR implies *less* warming for a fixed set of emissions, and our core results are robust to this modification.

Climate Damages. It is well-known that the damage function is consequential for estimating the social cost of carbon. Therefore, we consider three alternative specifications for damage functions, all of which are more pessimistic than the DICE damage function. Our first alternative allows for the economic effects of tipping points, following recent work by Dietz et al. (2021) and adopting their parameters. As a reminder, our baseline model uses the standard specification for damages in DICE2016, which is quadratic in temperature but includes no tipping points (see Equation 10). Our second alternative considers much larger damages than DICE, estimated in an influential paper by Burke et al. (2015). A third alternative considers the possibility that temperature also influences economic growth rates, as in Dell et al. (2012) and Moore and Diaz (2015). We calibrate parameters such

²⁰In fact, since writing this paper, a 2023 update to DICE was released that has incorporated key equations from FaIR in its updated climate module.

²¹We use an implementation of FaIR that was coded into the Julia programming language, where the rest of our model is run. Details are available at: <https://github.com/anthofflab/MimiFAIR.jl>.

that a 1-degree increase in temperature reduces GDP growth by 1 percentage point per year, consistent with the largest negative impacts on GDP growth presented by Moore and Diaz. For each of the three alternatives we summarize here, full details are provided in Appendix A.

Each of these damage function modifications substantially increase the economic costs of global warming under both population paths. Indeed, some of these model specifications have climate damages so severe that per capita output is lower in 2100 than it is today (see bottom left inset histogram in Figure 5). But because the differences in temperature are small across the two population paths, large damages per degree do not translate into large damage differences across population scenarios.

Population Emissions Intensity. In the baseline model, the industrial emissions come from economic production, not people, and so do not automatically scale identically with population. There is a lag between births and labor force increases. Therefore, an additional child today, who does not contribute to productive capacity, does not immediately increase emissions in the model.²²

To avoid the possibility of understating population’s effects on emissions, we mechanically increase the emissions elasticity of population to exactly one in each period. (See Appendix A). Overall, this does not change the quantity of emissions enough to make a qualitative difference to long-run warming.


Population → TFP Pass-Through. In the baseline model we calibrate λ, β in $g_{A,t} = \alpha_t L_t^\lambda A_t^{-\beta}$ to reflect leading empirical estimates (Bloom et al., 2020). We then calibrate α_t to replicate DICE’s TFP path.

To ensure that our findings do not rely on an overly optimistic calibration of how much more TFP a larger population can eventually generate, we make ideas “harder to find” by increasing β .²³ Specifically, we increase β from 3.1 to 4.5. This is a quantitatively meaningful change: TFP grows nearly 6-fold by 2200 with $\beta = 3.1$, but only doubles with $\beta = 4.5$ in our baseline model. Because it

²²Emissions increase once the child ages into the workforce. The consumption of a marginal individual who does not contribute to output is implicitly assumed to be substituting for some economic activity that would have otherwise taken place. This is consistent with the standard neoclassical structure of DICE and other medium- to long-run models of economic growth.

²³The long-run effect of a 1% increase in population is governed by the ratio of λ to β (see e.g., Eden and Kuruc, 2023), so this is quantitatively similar to scaling λ down by the same factor.

Figure 5: Net economic benefits of *Stabilization* are robust across 192 alternative sets of assumptions and model specifications, even though models vary widely



Notes: Alternative specifications are generated by crossing each of the six model dimensions indicated: climate policy (3 variants); climate modules (2 variants); climate damages (4 variants); amount of TFP growth (2 variants); source of TFP growth (2 variants); and emissions intensity of population (2 variants). The three inset histograms plot, for these 192 model specifications, the distributions of: year-2100 temperature change from pre-industrial under the *Depopulation* scenario (left); year-2200 temperature difference between *Stabilization* and *Depopulation* (right); and year-2100 GDP per capita under the *Depopulation* scenario (bottom). The histograms illustrate that these alternative models are substantially different, despite their convergent finding that net living standards are higher under *Stabilization* compared to *Depopulation*.

greatly shrinks the total growth of TFP, it closes the gap in TFP between population scenarios. However, even these smaller differences in TFP are enough for the larger population to eventually have higher per capita income net of climate damages.

Source of TFP Growth. So far, we have stressed the importance of people for the idea generation that produces TFP improvements. However, this is not always how endogenous growth models are specified. For example, Dietz and Stern (2015) implement the Romer (1986) endogenous growth model where economic capital is the key variable that scales innovation efforts. We implement a similar version of the innovation equation where TFP growth scales with total output—not people,

per se.

$$g_{A,t} = \alpha_t^Y Y_t^N A_t^{-\beta}. \quad (15)$$

Equation 15 recognizes that people need research labs, computers, and other productive economic capital to produce knowledge. Other things equal, a larger economy—meaning here the combination of people and other resources—can generate more new knowledge. Notice that this formulation is such that Y^N , output net of climate damages, determines growth. Therefore, like in Dietz and Stern (2015), climate damages influence growth rates indirectly by damaging the inputs to economic growth, making this similar to our third damage function modification. This ends up mattering very little to the main results.²⁴ Even a specification where capital was the *only* input to idea-creation would carry the implication that larger populations support larger capital stocks, which then support more TFP growth.

Figure 5 demonstrates that none of these modifications change the takeaways of the paper. Long-run per capita GDP is higher in the larger population for all 192 model combinations because the difference in warming is small in all 192 combinations.

5 Negative emissions generate an ambiguous relationship between population growth and climate outcomes

Our exercises in Sections 2 and 4 place a lower-bound on emissions at zero in order to isolate the potential increase in warming that population growth can generate. This section turns to a more novel possibility: the relationship between long-run levels of warming and near-term population growth rates may be *negative*. If scalable technologies come to exist by which GHGs can be removed from the atmosphere, a larger population will be able to produce more of this global public good. This is because removal of any fixed volume of atmospheric greenhouse gases will be a *fixed cost* from the perspective of a future generation. Therefore, for a larger population with a larger aggregate economy, the cost of removing a fixed

²⁴This is for two reasons. First, people are a primary input to Y^N , so net output is also substantially larger in *Stabilization* over the long run, due to the population increase. Second, capital in the economy increases with the size of the labor force. This is straightforward, as we are discussing aggregate capital, not per-worker capital.

volume of GHGs would be lower per capita and represent a smaller proportion of total output. Near-term population growth rate increases—leading to impacts on population levels only after a generations-long lag—may enable a more rapid reduction in the atmospheric GHG stock.

Futures with net-negative emissions are not particularly speculative. It is already understood how CO₂ can be captured from the atmosphere; whether these activities scale is primarily a matter of resource allocation. Indeed, the DICE model by default assumes that emissions eventually become net-negative. And DICE is not alone. Almost all of the pathways in a recent IPCC report on reaching 1.5°C include significant net-negative emissions in the second half of this century²⁵ and the United States’ first industrial direct air capture facility recently opened.²⁶ The uncertainty appears to be around *when*, not *whether*, the world will produce negative emissions at scale. If net annual per capita emissions become negative, then the long-run population-warming relationship will almost certainly become negative, over a long enough future.

To see this, consider a simple setting that relies on the following assumptions:

1. There is some t_{neg} after which negative emissions exceed positive emissions, independent of population size;
2. A larger population produces $\nu > 0$ tonnes more of these net-negative emissions per year; and
3. The larger population reaches time t_{neg} with $\Theta > 0$ more tonnes of CO₂ in the atmosphere.

In this setting, there will be some $t' > t_{neg}$ such that $\Theta = \nu(t' - t_{neg})$. This equality implies that the atmospheric stock of GHGs in the larger population future is equal to the stock in the smaller population future at time t' . Then, because the larger population continues to withdraw more GHGs beyond t' , the stock will be lower for all $t > t'$. Therefore, long-run warming will be lower in the larger population world.

²⁵See <https://www.iea.org/commentaries/going-carbon-negative-what-are-the-technology-options>.

²⁶See <https://www.nytimes.com/2023/11/09/climate/direct-air-capture-carbon.html>.

While this may be true in the limit, the argument does not tell us whether this result could arise on decision-relevant time horizons. To benchmark this, we use the timeline of net negative emissions in the original DICE model. We then compare the two population paths to evaluate how long after the introduction of net negative emissions it would be before the larger population would achieve the same level of warming as the smaller population.

In DICE's default settings, annual emissions become net-negative only beginning in 2150. After 2150 their quantity is governed by a cost-curve that maps incentives for negative emissions (i.e., a price on carbon), and a path for these prices. For this demonstration, the path of carbon prices is chosen such that the larger population future eventually generates -17 GtCO₂ annually, the maximum level of negative emissions that the baseline version of DICE produces prior to 2300. The same path of carbon prices is assumed in both *Stabilization* and *Depopulation*, in order to isolate the effect of population size.

Figure 6a plots the emissions paths of the two populations. Until 2150, this is an exact replication of the emissions paths in Figure 2. After 2150, more negative emissions are produced under the same policy incentives in *Stabilization* relative to *Depopulation* because the global economy is much larger by that time. Recall from Figure 1 that the population alone is more than 80% larger by 2200. The difference in output is greater still, because per capita productivity is greater in the stabilized world.²⁷ Just as a larger economy produces more non-rival bads, other things equal, it also produces more non-rival goods.


Panel (b) plots warming in the two population scenarios: In *Stabilization*, peak temperatures are slightly higher, exactly as in Figure 2, but the long-run temperature is substantially *lower* due to the increased resources for negative emissions. In this specification, it takes approximately 60 years from the point that negative emissions become possible to the point that the temperatures are equal.

Of course, the timing of net-negative emissions is subject to many uncertain details regarding future technologies and policies. And it will not, in general, be true that climate policy is invariant to population size. It is notable, though, that under


²⁷The larger population could produce more negative emissions even if we eliminated the productivity benefits of population growth and assumed the two populations were equally wealthy on a per capita level. The first-order issue is one of scale.

Figure 6: Larger populations can have climate *benefits* if emissions become negative

(a) A larger population produces more positive and negative emissions...



(b) ...which can reverse the initial temperature increase



Notes: A replication of the exercise that determines emissions and temperature effects in Figure 2, but that allows emissions to become net-negative after 2150 (the year in which the 2016 revision of DICE assumes negative emissions begin). In the main analysis and all specifications in Figure 5, we constrain the model such that annual GHG emissions can never fall below zero. The plots illustrate that with negative emissions technologies, there is the potential for climate *benefits* of a larger long-run population.

this set of accessible projections and assumptions provided by DICE, the time it takes for the larger population to have better climate outcomes is on the order of 50, rather than 500, years from the time net-negative emissions are realized. Moreover, the qualitative point will hold regardless of the details: population growth influences long-run populations, and humanity's long-run climate challenge is likely to be collectively financing the removal of existing emissions.

6 Discussion and Conclusion

Global fertility is unprecedentedly low and is continuing to fall. This has prompted concern over an aging and shrinking workforce, but also optimism about environmental benefits. Foremost among the supposed environmental benefits are reduced greenhouse gas emissions and lower levels of long-run warming. This paper shows that this optimism greatly overstates the potential climate benefits of further declines in fertility: Feasible emissions reductions resulting from changes in population dynamics are small when compared against well-studied productivity benefits of near-term increases in population growth. So population reduction is not a substitute for decarbonization policies. Moreover, reductions in fertility do not *complement* other efforts towards decarbonization. Decarbonization efforts reduce per capita emissions, reducing the marginal impact of declining populations and rendering future population sizes less important—or altogether unimportant, once net emissions per person reach zero.

There are at least three ways that future research could push beyond our paper. First, although climate change is most widely discussed potential environment impact of population size, it is not the only one. A larger population would have environmental effects beyond climate change, including on biodiversity, non-human animals, and non-carbon air and water pollution. Our main analysis does not aim to address these. Additionally, we do not account for the possibility that it would be less costly for a population to invest in the human capital of a smaller generation of children, as is described in the macroeconomics of fertility literature (see e.g., Galor, 2022).

Second, productivity growth is not the only potential welfare benefit of a larger population. In our analysis, we present results in per capita terms, and give no advantage to a larger population future for the reason that more people get to exist. Ignoring this pathway ignores potentially one of the major social welfare benefits of increased fertility (see e.g., Klenow et al., 2023).

Third, our analysis also omits regional heterogeneity. We consider a marginal increase in *global* fertility rates. But population growth is projected to be geographically uneven in the coming decades. For example, most future population growth will take place in Africa. That would be true in *Stabilization* and *Depopulation* and

does not imply anything about where the marginal population growth we study comes from. Our exercise is one that imagines a representative agent of the global population having a slightly larger family to isolate the effect of population size. The core mechanism—that population size changes are slow and predictable relative to the urgency of emissions reduction—is a fact that will apply at finer levels of geographic aggregation. Like the other modeling variants examined in Figure 5, substituting a regional model would not alter these core results.

In general, as societies prepare for the concurrent challenges of reducing emissions and confronting sustained low fertility, it is useful to understand that these objectives are not at odds in a quantitatively meaningful way.

References

- Arkolakis, Costas and Conor Walsh (2023) "Clean Growth," *NBER Working Paper No. 31615*.
- Basten, Stuart, Wolfgang Lutz, and Sergei Scherbov (2013) "Very long range global population scenarios to 2300 and the implications of sustained low fertility," *Demographic Research*, 28, 1145–1166.
- Bloom, Nicholas, Charles I Jones, John Van Reenen, and Michael Webb (2020) "Are Ideas Getting Harder to Find?" *American Economic Review*, 110 (4), 1104–44.
- Bongaarts, John and Brian C O'Neill (2018) "Global warming policy: Is population left out in the cold?" *Science*, 361 (6403), 650–652.
- Bradshaw, Corey JA and Barry W Brook (2014) "Human population reduction is not a quick fix for environmental problems," *Proceedings of the National Academy of Sciences*, 111 (46), 16610–16615.
- Bretschger, Lucas (2020) "Malthus in the light of climate change," *European Economic Review*, 127, 103477.
- Budolfson, Mark and Dean Spears (2021) "Population ethics and the prospects for fertility policy as climate mitigation policy," *Journal of Development Studies*, 1–12.
- Burke, Marshall, Solomon M Hsiang, and Edward Miguel (2015) "Global non-linear effect of temperature on economic production," *Nature*, 527 (7577), 235–239.
- Casey, Gregory and Oded Galor (2017) "Is faster economic growth compatible with reductions in carbon emissions? The role of diminished population growth," *Environmental Research Letters*, 12 (1), 10–1088.
- Cruz, José-Luis and Esteban Rossi-Hansberg (Forthcoming) "The Economic Geography of Global Warming," *Review of Economic Studies*.
- Dasgupta, Partha, Aisha Dasgupta, and Scott Barrett (2021) "Population, Ecological Footprint and the Sustainable Development Goals," *Environmental and Resource Economics*, 1–17.

- Dell, Melissa, Benjamin F Jones, and Benjamin A Olken (2012) "Temperature shocks and economic growth: Evidence from the last half century," *American Economic Journal: Macroeconomics*, 4 (3), 66–95.
- Dietz, Simon, James Rising, Thomas Stoerk, and Gernot Wagner (2021) "Economic impacts of tipping points in the climate system," *Proceedings of the National Academy of Sciences*, 118 (34).
- Dietz, Simon and Nicholas Stern (2015) "Endogenous growth, convexity of damage and climate risk: how Nordhaus' framework supports deep cuts in carbon emissions," *Economic Journal*, 125 (583), 574–620.
- Eden, Maya and Kevin Kuruc (2023) "The long-run relationship between per capita incomes and population sizes," *CEPR Discussion Paper DP19353*.
- Ekerdt, Lorenz and Kai-Jie Wu (2023) "Self-Selection and the Diminishing Returns to Research," *Working Paper*, Available at <https://sites.google.com/view/lorenz-ekerdt/research>.
- Fernández-Villaverde, Jesús, Gustavo Ventura, and Wen Yao (2023) "The Wealth of Working Nations," *NBER Working Paper No. 31914*.
- Galor, Oded (2022) *The Journey of Humanity: The origins of wealth and inequality*: Penguin.
- Gerlagh, Reyer, Veronica Lupi, and Marzio Galeotti (2023) "Fertility and climate change," *The Scandinavian Journal of Economics*, 125 (1), 208–252.
- Hausfather, Zeke and Glen P Peters (2020) "Emissions—the 'business as usual' story is misleading," *Nature*, 577 (7792), 618–620.
- Henderson, J Vernon, Bo Yeon Jang, Adam Storeygard, and David N. Weil (2023) "Climate Change, Population Growth, and Population Pressure," *Working Paper*.
- Hopenhayn, Hugo, Julian Neira, and Rish Singhania (2022) "From population growth to firm demographics: Implications for concentration, entrepreneurship and the labor share," *Econometrica*, 90 (4), 1879–1914.

- Howard, Peter H and Thomas Sterner (2017) "Few and not so far between: a meta-analysis of climate damage estimates," *Environmental and Resource Economics*, 68 (1), 197–225.
- Jones, Charles I (1995) "R&D-Based Models of Economics Growth," *Journal of Political Economy*, 103 (4).
- (2022) "The end of economic growth? Unintended consequences of a declining population," *American Economic Review*, 112 (11), 3489–3527.
- Karahan, Fatih, Benjamin Pugsley, and Ayşegül Şahin (2019) "Demographic origins of the startup deficit," *National Bureau of Economic Research Working Paper*, No. 25874.
- Klenow, Peter, Charles I Jones, Mark Bils, and Mohamad Adhami (2023) "Population and Welfare: The Greatest Good for the Greatest Number," Available at: <http://web.stanford.edu/~chadj/popwelfare.pdf>.
- Kruse-Andersen, Peter (2023a) "Directed technical change, environmental sustainability, and population growth," *Journal of Environmental Economics and Management*, 122.
- (2023b) "Testing R&D-Based Endogenous Growth Models," *Oxford Bulletin of Economics and Statistics*, 85 (5), 1083–1110.
- Kuruc, Kevin, Sangita Vyas, Mark Budolfson, Michael Geruso, and Dean Spears (2023) "Population Decline: Too Small and Too Slow to Influence Climate Change," *Population Wellbeing Initiative Working Paper*, Available at: https://nbviewer.org/github/kevinkuruc/Papers_for_Download/blob/main/VYPR_Resubmission.pdf.
- Maestas, Nicole, Kathleen J Mullen, and David Powell (2023) "The Effect of Population Aging on Economic Growth, the Labor Force, and Productivity," *American Economic Journal: Macroeconomics*, 15 (2), 306–332.
- Marois, Guillaume, Stuart Gietel-Basten, and Wolfgang Lutz (2021) "China's low fertility may not hinder future prosperity," *Proceedings of the National Academy of Sciences*, 118 (40).

- Moore, Frances C and Delavane B Diaz (2015) “Temperature impacts on economic growth warrant stringent mitigation policy,” *Nature Climate Change*, 5 (2), 127–131.
- National Academies (2017) “Valuing Climate Damages: Updating Estimation of the Social Cost of Carbon Dioxide,” report, National Academies of Sciences, Engineering, and Medicine.
- Nordhaus, William D (2017) “Revisiting the social cost of carbon,” *Proceedings of the National Academy of Sciences*, 114 (7), 1518–1523.
- O’Neill, Brian C, Brant Liddle, Leiwen Jiang, Kirk R Smith, Shonali Pachauri, Michael Dalton, and Regina Fuchs (2012) “Demographic change and carbon dioxide emissions,” *Lancet*, 380 (9837), 157–164.
- Ou, Yang, Gokul Iyer, Leon Clarke et al. (2021) “Can updated climate pledges limit warming well below 2° C?” *Science*, 374 (6568), 693–695.
- Our World in Data (2023) “Birth Rates,” <https://ourworldindata.org/grapher/crude-birth-rate?tab=chart>, Accessed: December 27, 2023.
- Peters, Michael (2022) “Market size and spatial growth—Evidence from Germany’s post-war population expulsions,” *Econometrica*, 90 (5), 2357–2396.
- Peters, Michael and Conor Walsh (2021) “Population Growth and Firm Dynamics,” *National Bureau of Economic Research Working Paper*, No. 29424.
- Pindyck, Robert S (2022) “Population, Productivity, and Sustainable Consumption,” *NBER Working Paper*, No. 30019.
- Raftery, Adrian E and Hana Ševčíková (2023) “Probabilistic population forecasting: Short to very long-term,” *International journal of forecasting*, 39 (1), 73–97.
- Romer, Paul M (1986) “Increasing Returns and Long-Run Growth,” *Journal of Political Economy*, 94 (5), 1002–1037.
- (1990) “Endogenous Technological Change,” *Journal of Political Economy*, 98 (5, Part 2), S71–S102.

Sherwood, Steven C, Alex Sen Gupta, and Stephen E Schwartz (2022) “Probability of committed warming exceeding 1.5° C and 2.0° C Paris targets,” *Environmental Research Letters*, 17 (6), 064022.

Spears, Dean, Sangita Vyas, Gage Weston, and Mike Geruso (2023) “Long-term population projections: Scenarios of low or rebounding fertility,” *PWI Working Paper No. 2306*, Available at: https://sites.utexas.edu/pwi/files/2023/09/spike_site.pdf.

Stone, Lyman (2020) “Pro-Natal Policies Work, But They Come With a Hefty Price Tag,” *Institute for Family Studies, Blog Post*. Available at: <https://ifstudies.org/blog/pro-natal-policies-work-but-they-come-with-a-hefty-price-tag>.

United Nations (2022) “World population prospects 2022,” *New York, NY: United Nations Department for Economic and Social Affairs*.

Vollrath, Dietrich (2020) *Fully grown*: University of Chicago Press.

Weil, David (2023) “Replacement Fertility is Neither Natural nor Optimal nor Likely,” *Working Paper*, Available at: https://conference.nber.org/conf_papers/f179025.pdf.

Weil, David N (1999) “Population Growth, Dependency, and Consumption,” *American Economic Review*, 89 (2), 251–255.

Wynes, Seth and Kimberly A Nicholas (2017) “The climate mitigation gap: education and government recommendations miss the most effective individual actions,” *Environmental Research Letters*, 12 (7), 074024.

Zhou, Anson (2023) “Bounding fertility elasticities,” *Economics Letters*, 228, 111151.

Appendix

A Details of modifications in Section 4.2

Damage Functions

Dietz et al. present a reduced-form, additive modification of standard quadratic damage functions with coefficients ξ_1, ξ_2 . Note that a damage function with tipping points can still take a smooth quadratic form because there is *ex-ante* uncertainty about where such tipping points lie. Increases in temperature can continuously increase the probability of reaching a tipping point, even if the *ex-post* damage function has a discontinuity.

$$D_t = (\psi_1 + \xi_1)T_t + (\psi_2 + \xi_2)T_t^2$$

We use the coefficients reported in Figure 5 of Dietz et al. (2021).

Our second alternative considers much larger damages than DICE, estimated in an influential paper by Burke et al. (2015). These damage estimates come from a non-linear model disaggregated to the country level. DICE is specified at a coarser level of aggregation, so we implement the reduced-form version presented in Figure 5d of Burke et al., linking global temperatures to global losses of GDP. We translate the graphical depiction to numerical values using data extraction software and estimate a cubic function, $D = \alpha_1 T + \alpha_2 T^2 + \alpha_3 T^3$, for the corresponding damage function.

Our third alternative considers the possibility that temperature also influences economic growth rates, as in Dell et al. (2012) and Moore and Diaz (2015). Moore and Diaz (2015) implement this in a regional integrated assessment model, making exact replication infeasible in our global setting. We instead implement their functional form at the global level and employ coefficients on the higher end of their proposed range in an effort to be conservative (against our findings). Specifically, the rate of TFP growth becomes:

$$g_{A,t} = \alpha_t L_t^\lambda A_t^{-\beta} - \varepsilon \tilde{T}.$$

We calibrate ε such that a 1-degree increase in \tilde{T} reduces GDP growth by 1 percentage point per year, consistent with the largest negative impacts on GDP growth presented by Moore and Diaz. Also following their implementation, we use what

they call “effective temperature,” \tilde{T} , to allow for adaptation. Subtracting a function of past temperatures allows for the long-run effect of a fixed level of warming to be reduced over time, in their specification back to zero. Our numerical implementation is slightly different than Diaz and Moore owing to differences in model construction, but we retain that warming (i) passes through to \tilde{T} one-to-one in the immediate-term and (ii) has a near-zero effect on growth rates after 30 years at that level. Specifically, Moore and Diaz define $\tilde{T}_t = \sum_{j=1850}^{j=t} (T_j - T_{j-1}) e^{-a(t-j)}$ such that if warming is fixed at some level in the long-run, the $\tilde{T} \rightarrow 0$. In other words, there is no effect in the long-run of one-time increase in global temperatures. For simplicity, we instead subtract a rolling average of the prior 30 years. In Moore and Diaz’s calibration, \tilde{T} is nearly zero after 30 years following a one-time shock, so our simplification captures similar transition dynamics.

Population Emissions Elasticity


Beyond scaling land-use emissions to population as we do in every model interaction, to avoid understating the effect of population on emissions, we redefine industrial emissions as

$$E_{Ind,t} = (1 - \mu_t) \times \sigma_t^N \times N_t,$$

where σ^N corresponds to emissions per capita, rather than emissions per unit of output (also recall that μ is the mitigation rate). This functional form ensures that if in period t *Stabilization* has a population 10% larger than *Depopulation*, emissions will also be 10% larger. We calibrate σ^N to again replicate DICE2016’s baseline outcomes; i.e., we fit the equation $\sigma_t^N \times N_{t,DICE} = \sigma_t \times Y_{t,DICE}$ for DICE’s population and output. This prevents this redefinition from substantively changing anything about the baseline cases we build from.


Supplementary Figures for Online Publication

Figure A1: Our modified model reproduces DICE's output with DICE's population



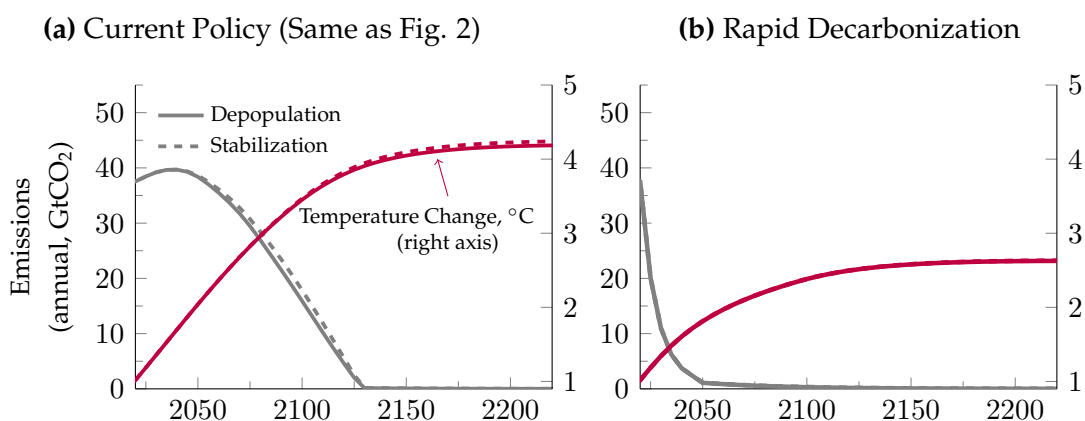
Notes: Verification that the modified version of DICE—with endogenized TFP and land-use emissions—exactly replicates DICE2016R when the original DICE population and policy trajectory is assumed. The output from DICE2016R is available at <https://williamnordhaus.com/dicerice-models>.

Figure A2: Industrial emissions in Figure 2 are calibrated to independent projections



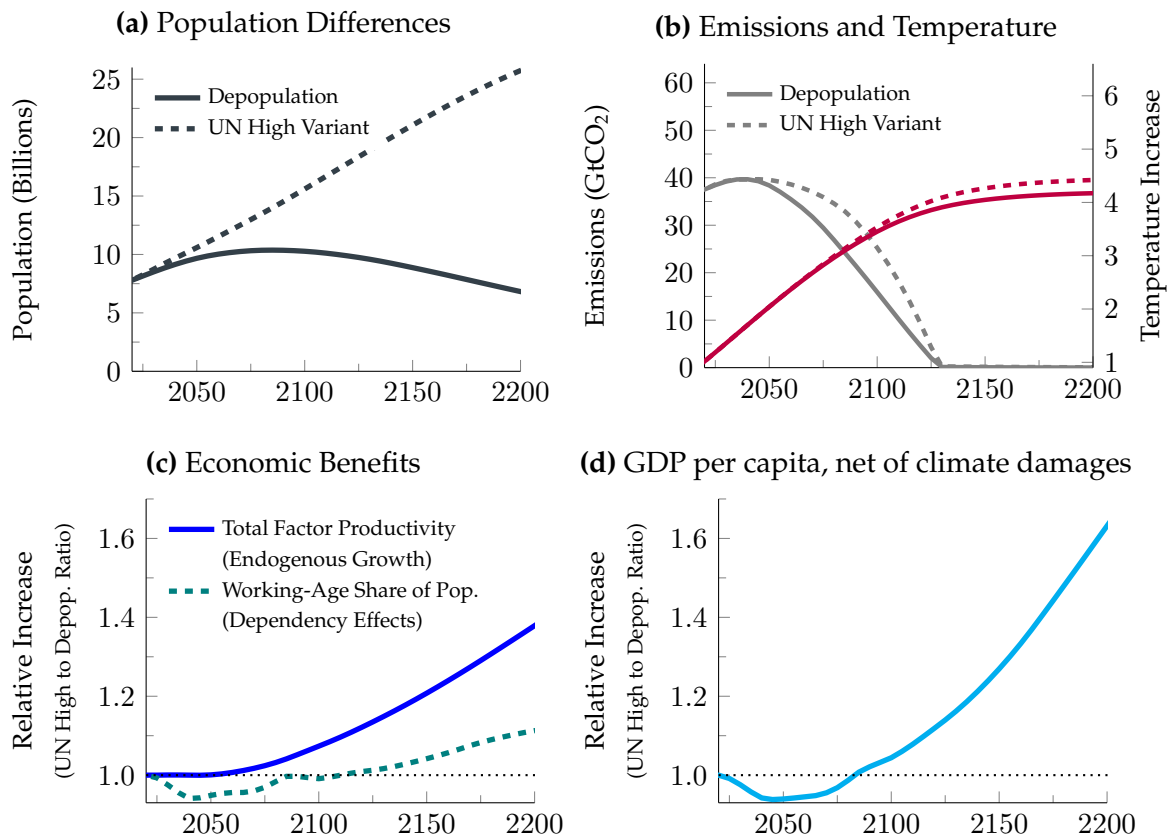
Notes: The industrial emissions pathways in our model plotted against the projected industrial emissions in Ou et al. (2021) (from their Table S8). To calibrate our emissions pathway, we use a single variable—a fixed z p.p. increase in mitigation rates, μ_t , each five-year period (until an upper bound of 100% mitigation). DICE assumes a background rate of technological progress (σ_t) that increases over time, providing the accelerating decline towards the end of the century. With $z = 4$ p.p. per five-year time period, emissions initially decline less quickly in our model than is projected by Ou et al. (2021). By about 2080, the decline in our scenario is slightly more rapid. Note that the shape of an accelerating decline is consistent with other leading projections (see e.g., Hausfather and Peters, 2020, who uses International Energy Association projections with a similar time-profile to our path).

Figure A3: Depiction of optimistic policy scenario used in Figure 5



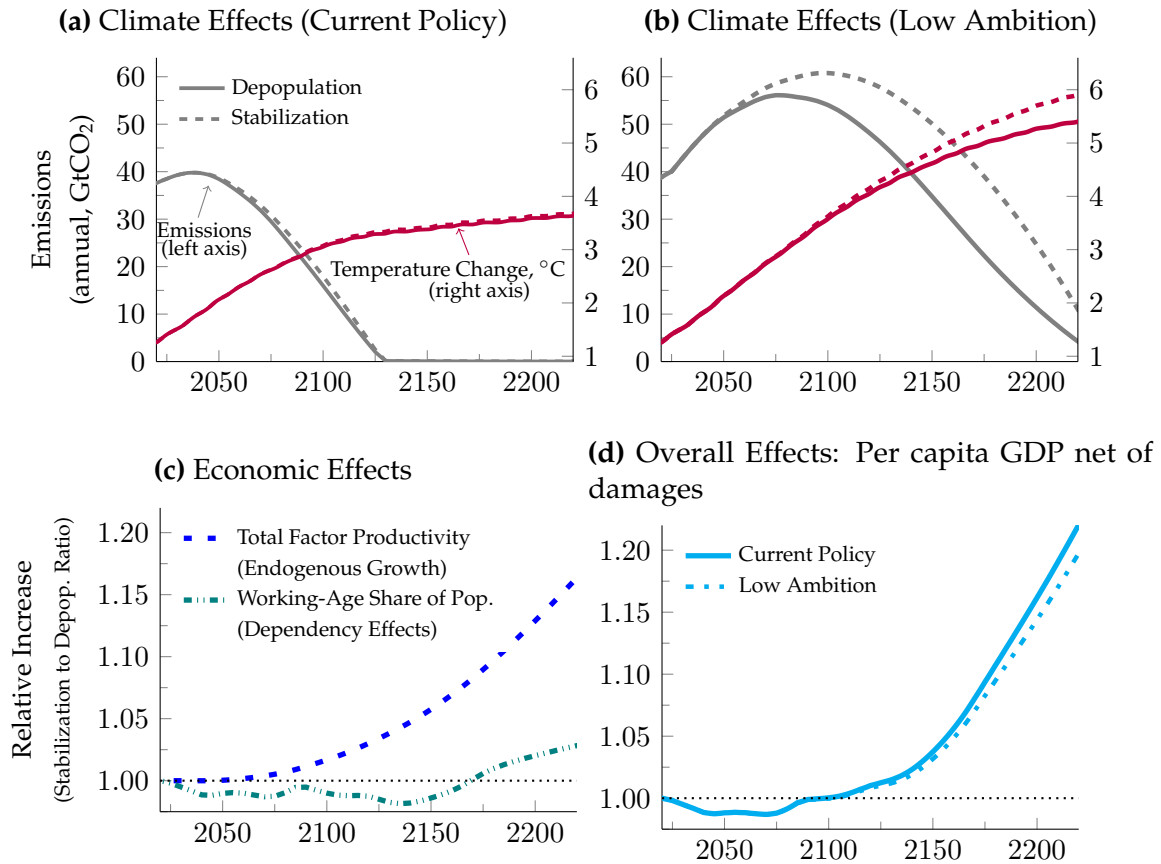
Notes: Emissions and warming under our “ambitious” climate policy, where emissions fall close to net-zero by around 2050. Current policy plotted in left panel for reference. This panel is identical the panel (a) in Figure 3. This scenario is extremely ambitious relative to projection of future emissions. The point is to formally demonstrate that *if* emissions reductions are rapid, there is no distinction in warming between larger and smaller populations. As a corollary for policymakers and others interested in achieving rapid decarbonization: population change will not make any difference in whether ambitious temperature targets are met.

Figure A4: Benefits of population remain large when comparing UN High population projection with *Depopulation*



Notes: A replication of our results where an extension of the UN High variant, rather than *Stabilization*, is compared with *Depopulation*. Uses “current policy” mitigation pathways in both population scenarios.

Figure A5: Replication of our results using FaIR climate module



Notes: A replication of our results in which the climate representation has been replaced by the FaIR model. FaIR produces less warming conditional on the same path of emissions (see Figure 2 where warming exceeds 4°C for this emissions path with DICE's atmospheric module). However, the *difference* in warming between population scenarios remains small, and therefore *Stabilization* continues to have a significant advantage in long-run GDP per capita relative to *Depopulation*.